

ProvBench-RAG: A Benchmark Resource for Provenance-Sensitive Retrieval Under Near-Duplicate Source Confusion

[Authors — Resource track, single-blind]

Track: CIKM 2026 Resource

Keywords: benchmark, retrieval, provenance, RAG, dataset, evaluation

Abstract

Retrieval-augmented systems are often evaluated on whether some passage supports an answer, but deployed pipelines must also choose the correct document instance among near-duplicates (mirrors, stale snapshots, scraped copies, jurisdiction variants). We release ProvBench-RAG, an open benchmark resource comprising JSONL schemas, three reproducible tracks (1,500 synthetic clusters, 140 Wikipedia revision clusters, 1,640 pooled queries), deterministic evaluation code, and reference baseline rankings. The resource exposes preferred-source accuracy (PSA@1) and composite ProvScore alongside recall and nDCG. On the synthetic track, lexical BM25 achieves PSA@1 of 0 under planted decoys while a bundled provenance reranker reaches 0.75; Wikipedia and pooled tracks quantify transfer and heterogeneous-index stress. All data, metrics, and paper tables regenerate from public scripts and frozen annotations in the repository.